

CS 3551 DISTRIBUTED COMPUTING

UNIT I INTRODUCTION

8

Introduction: Definition-Relation to Computer System Components – Motivation – Message - Passing Systems versus Shared Memory Systems – Primitives for Distributed Communication – Synchronous versus Asynchronous Executions – Design Issues and Challenges; A Model of Distributed Computations: A Distributed Program – A Model of Distributed Executions – Models of Communication Networks – Global State of a Distributed System.

LIKE



COMMENT




SHARE




SUBSCRIBE



Distributed Program




- A distributed program is composed of a set of n asynchronous processes $p_1, p_2, \dots, p_i, p_n$ that communicate by message passing over the communication network.
- We assume that each process is running on a different processor.
- The processes do not share a global memory and communicate solely by passing messages.
- C_{ij} - Channel from p_i to process p_j
- m_{ij} - a message sent by p_i to p_j .
- Don't share a global clock. ✓
- Process execution and message transfer are asynchronous
- The global state of a distributed computation is composed of the states of the processes and the communication channels
- The state of a process is characterized by the state of its local memory and depends upon the context. The state of a channel is characterized by the set of messages in transit in the channel



Model for Distributed Execution

• -internal

no
process id



binary
relation
 $H_i = (h_i, \vec{i})$
↓
events

⊗ atomic

- 1) internal \rightarrow process
- 2) send \rightarrow sender + channel
- 3) receive \rightarrow receiver + channel

Model for Distributed Execution



1. Execution of a process consists of a sequential execution of its actions.
2. The actions are atomic and the actions of a process are modeled as three types of events, internal events, message send events(send (m)), and message receive events(rec(m)).
3. The occurrence of events changes the states of respective processes and channels, thus causing transitions in the global system state.
4. An internal event changes the state of the process at which it occurs.
5. A send event (or a receive event) changes the state of the process that sends (or receives) the message and the state of the channel on which the message is sent (or received)

The events at a process are linearly ordered by their order of occurrence. The execution of process p_i produces a sequence of events $e_i^1, e_i^2, \dots, e_i^x, e_i^{x+1}, \dots$ and is denoted by \mathcal{H}_i :

$$\mathcal{H}_i = (h_i, \rightarrow_i),$$

$$e_i^1 \rightarrow e_i^2$$

where h_i is the set of events produced by p_i and binary relation \rightarrow_i defines a linear order on these events. Relation \rightarrow_i expresses causal dependencies among the events of p_i .

The send and the receive events signify the flow of information between processes and establish causal dependency from the sender process to the receiver process. A relation \rightarrow_{msg} that captures the causal dependency due to message exchange, is defined as follows. For every message m that is exchanged between two processes, we have


$$send(m) \rightarrow_{msg} rec(m).$$

Relation \rightarrow_{msg} defines causal dependencies between the pairs of corresponding send and receive events.

Causal Precedence Relation

① $e_i \rightarrow e_j$
 $e_j \rightarrow e_i$
 $e_j || e_i$

$$e_i \rightarrow e_j$$



x_2
 y
 $e_1^1 \rightarrow e_3$
 v_1

- ① LHS \rightarrow occur first
- ② LHS info \Rightarrow also available for RHS

$$e_i \nrightarrow e_j$$

Rules

- ① • for any two events e_i and e_j , $e_i \nrightarrow e_j \not\Rightarrow e_j \nrightarrow e_i$
- ② • for any two events e_i and e_j , $e_i \rightarrow e_j \Rightarrow e_j \nrightarrow e_i$.

For any two events e_i and e_j , if $e_i \nrightarrow e_j$ and $e_j \nrightarrow e_i$, then events e_i and e_j are said to be concurrent and the relation is denoted as $e_i || e_j$.

Causal precedence relation

The execution of a distributed application results in a set of distributed events produced by the processes. Let $H = \cup_i h_i$ denote the set of events executed in a distributed computation. Next, we define a binary relation on the set H , denoted as \rightarrow , that expresses causal dependencies between events in the distributed execution.

$$\forall e_i^x, \forall e_j^y \in H, \underbrace{e_i^x \rightarrow e_j^y} \Leftrightarrow \left\{ \begin{array}{l} e_i^x \rightarrow_i e_j^y \text{ i.e., } (i = j) \wedge (x < y) \\ \text{or} \\ e_i^x \rightarrow_{msg} e_j^y \\ \text{or} \\ \exists e_k^z \in H : e_i^x \rightarrow e_k^z \wedge e_k^z \rightarrow e_j^y \end{array} \right.$$

The causal precedence relation induces an irreflexive partial order on the events of a distributed computation [6] that is denoted as $\mathcal{H} = (H, \rightarrow)$.

Logical vs Physical Concurrency

$$\frac{e_1}{e_3}$$

- In a distributed computation, two events are logically concurrent if and only if they do not causally affect each other.
- Physical concurrency, on the other hand, has a connotation that the events occur at the same instant in physical time.
- Note that two or more events may be logically concurrent even though they do not occur at the same instant in physical time.

$$\Rightarrow \{e_1^3, e_2^4, e_3^3\}$$



LIKE



COMMENT



SHARE




SUBSCRIBE



Models of Communication Networks (FIFO, Non-FIFO, Causal Ordering)

1. **FIFO** - each channel acts as a first-in first-out message queue and thus, message ordering is preserved by a channel.
2. **non-FIFO model**, a channel acts like a set in which the sender process adds messages and the receiver process removes messages from it in a random order.
3. **Causal Ordering (built-in synch)**- based on Lamport's "happens before" relation. A system that supports the causal ordering model satisfies the following property:

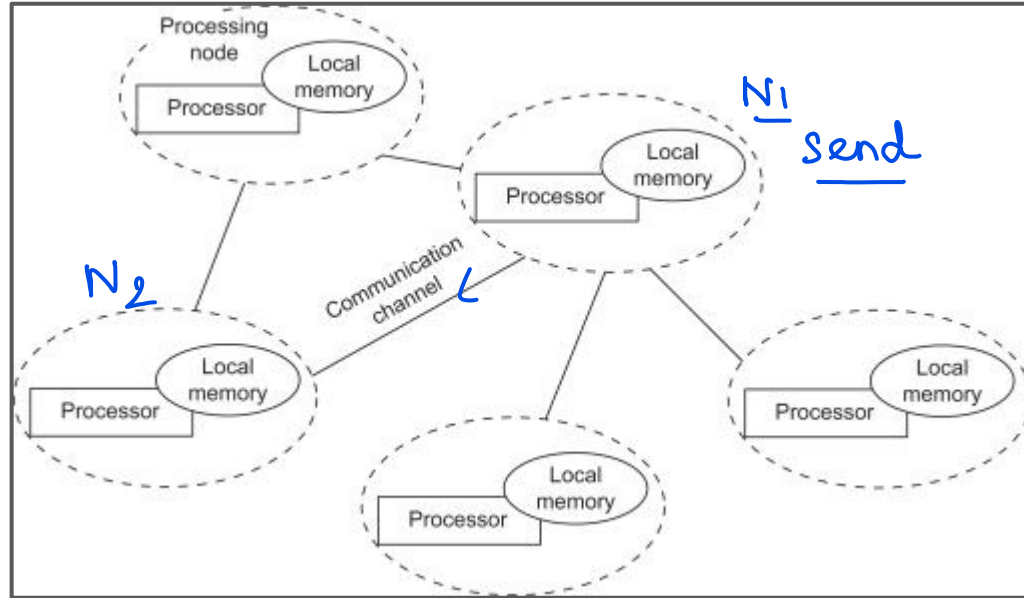


CO: For any two messages m_{ij} and m_{kj} , if $send(m_{ij}) \rightarrow send(m_{kj})$,
then $rec(m_{ij}) \rightarrow rec(m_{kj})$.

That is, this property ensures that causally related messages destined to the same destination are delivered in an order that is consistent with their causality relation. Causally ordered delivery of messages implies FIFO message delivery. Furthermore, note that $CO \subset FIFO \subset Non-FIFO$.

Global State of Distributed System

process
channels



1. The state of a process at any time is defined by the contents of processor registers, stacks, local memory, etc. and depends on the local context of the distributed application.
2. The state of a channel is given by the set of messages in transit in the channel.
3. The occurrence of events changes the states of respective processes and channels, thus causing transitions in global system state.
 - a. For example, an internal event changes the state of the process at which it occurs. A send event (or a receive event) changes the state of the process that sends (or receives) the message and the state of the channel on which the message is sent (or received).

State of Process


$$p_i \Rightarrow LS_i^x \Rightarrow \underset{\text{process}}{LS_i^1} \rightarrow \text{event: complete}$$

$$p_i \rightarrow e_i^1, e_i^2, e_i^3$$

Let LS_i^x denote the state of process p_i after the occurrence of event e_i^x and before the event e_i^{x+1} . LS_i^0 denotes the initial state of process p_i .

① $send(m) \leq LS_i^x$ denote the fact that $\exists y: 1 \leq y \leq x :: e_i^y = send(m)$
 $rec(m) \not\leq LS_i^x$ denote the fact that $\forall y: 1 \leq y \leq x :: e_i^y \neq rec(m)$.

State of Channel



Let $SC_{ij}^{x,y}$ denote the state of a channel C_{ij} defined as follows.


$$SC_{ij}^{x,y} = \{m_{ij} \mid \text{send}(m_{ij}) \leq LS_i^x \wedge rec(m_{ij}) \not\leq LS_j^y\}.$$

Thus, channel state $SC_{ij}^{x,y}$ denotes all messages that p_i sent up to event e_i^x and which process p_j had not received until event e_j^y .

How to find global state?

$$GS = \{\bigcup_i LS_i^{x_i}, \bigcup_{j,k} SC_{jk}^{y_j, z_k}\}$$


- ❖ For a global snapshot to be meaningful, the states of all the components of the distributed system must be recorded at the same instant.
- ❖ This will be possible if the local clocks at processes were perfectly synchronized or there was a global system clock that could be instantaneously read by the processes. However, both are impossible.
- ❖ Solution: ✓
- ❖ Recording at different time will be meaningful provided every message that is recorded as received is also recorded as sent.
- ❖ **Basic idea is that an effect should not be present without its cause.**
- ❖ **States that don't violate causality are called consistent global states and are meaningful global states**




A global state $GS = \{\bigcup_i LS_i^{x_i}, \bigcup_{j,k} SC_{jk}^{y_j, z_k}\}$ is a *consistent global state* iff it satisfies the following condition:

$$\forall m_{ij} : send(m_{ij}) \not\leq LS_i^{x_i} \Rightarrow m_{ij} \notin SC_{ij}^{x_i, y_j} \wedge rec(m_{ij}) \not\leq LS_j^{y_j}$$


That is, channel state $SC_{ik}^{y_i, z_k}$ and process state $LS_k^{z_k}$ must not include any message that process p_i sent after executing event $e_i^{x_i}$.





Inconsistent state

of local states $\{LS_1^1, LS_2^3, LS_3^3, LS_4^2\}$ is inconsistent because the state of p_2 has recorded the receipt of message m_{12} , however, the state of p_1 has not recorded its send.



Consistent state

states $\{LS_1^2, LS_2^4, LS_3^4, LS_4^2\}$ is consistent; all the channels are empty except C_{21} that contains message m_{21} .

Strongly Consistent ✓

A global state $GS = \{\bigcup_i LS_i^{x_i}, \bigcup_{j,k} SC_{jk}^{y_j, z_k}\}$ is transitless iff

$$\forall i, \forall j : 1 \leq i, j \leq n :: SC_{ij}^{y_i, z_j} = \phi.$$

Thus, all channels are recorded as empty in a transitless global state. A global state is *strongly consistent* iff it is transitless as well as consistent. Note that in Figure 2.2, the global state consisting of local states $\{LS_1^2, LS_2^3, LS_3^4, LS_4^2\}$ is strongly consistent.



LIKE



COMMENT



SHARE



SUBSCRIBE

